

01 – EIS & SOC Exploration

Chouaib (github.com/ChouaibB)

December 19, 2025

Contents

1. Data source: SoC EIS LFP dataset	2
2. Load data & quick sanity checks.	2
2. Cycle QC and averaging	4
3. EIS visualisation vs SOC	5
4. Feature engineering: simple impedance fingerprints	7
5. SOC prediction target & problem setup	12
6. Summary of SOC modelling experiment	16

1. Data source: SoC EIS LFP dataset

For this notebook I use the public dataset:

Mustafa, Hamza; Bourelly, Carmine; Vitelli, Michele; Milano, Fillippo; Molinara, Mario; Ferrigno, Luigi (2024),

“SoC Estimation on Li-ion Batteries: A New EIS-based Dataset for data-driven applications”,

Mendeley Data, V2, doi: 10.17632/cb887gkmxw.2.

Available at: <https://data.mendeley.com/datasets/cb887gkmxw/2>

License: **CC BY 4.0**

Dataset highlights:

- 11 cylindrical **LFP** cells (B01–B11).
- EIS at 20 SoC levels: 100%, 95%, 90%, 85%, 80%, 75%, 70%, 65%, 60%, 55%, 50%, 45%, 40%, 35%, 30%, 25%, 20%, 15%, 10%, 5%.
- Two discharge cycles per cell, same frequency grid from 0.01 Hz to 1000 Hz (28 frequencies).

Folder structure (after downloading the dataset into `../data/soc_eis_lfp/`):

- One folder per cell: B01, B02, ..., B11.
- Inside each Bxx folder:
 - **EIS Measurement/** — subfolders with the actual EIS CSV files.
 - **Capacity Measurement/** — capacity test CSV (not used in this first notebook).

2. Load data & quick sanity checks.

```
Shape: (12320, 7)
Batteries: ['B01' 'B02' 'B03' 'B04' 'B05' 'B06' 'B07' 'B08' 'B09' 'B10' 'B11']
Cycles: [np.int64(1), np.int64(2)]
Unique SOCs: [np.int64(5), np.int64(10), np.int64(15), np.int64(20),
np.int64(25), np.int64(30), np.int64(35), np.int64(40), np.int64(45),
np.int64(50), np.int64(55), np.int64(60), np.int64(65), np.int64(70),
np.int64(75), np.int64(80), np.int64(85), np.int64(90), np.int64(95),
np.int64(100)]
Freq min/max: 0.01 → 1000.0
Temp min/max: 20 → 20
```

```
      frequency_hz  z_real_ohm  z_imag_ohm  temperature_c battery_id  cycle  soc
0            0.01    0.227844   -0.313990           20        B01      1  100
1            0.02    0.182858   -0.197273           20        B01      1  100
2            0.03    0.163270   -0.160068           20        B01      1  100
3            0.05    0.138709   -0.113159           20        B01      1  100
4            0.08    0.121685   -0.082724           20        B01      1  100

      frequency_hz  z_real_ohm  z_imag_ohm  temperature_c battery_id \
count       12320.000  12320.000  12320.000     12320.0          12320
unique        NaN        NaN        NaN          NaN            NaN          11
top          NaN        NaN        NaN          NaN            NaN          B01
freq          NaN        NaN        NaN          NaN            NaN         1120
```

```
mean      117.360      0.078     -0.015      20.0      NaN
std       246.764      0.025      0.031      0.0      NaN
min       0.010      0.042     -1.329     20.0      NaN
25%       0.275      0.061     -0.015     20.0      NaN
50%       6.500      0.072     -0.007     20.0      NaN
75%      85.750      0.089     -0.006     20.0      NaN
max      1000.000      0.389     -0.001     20.0      NaN
```

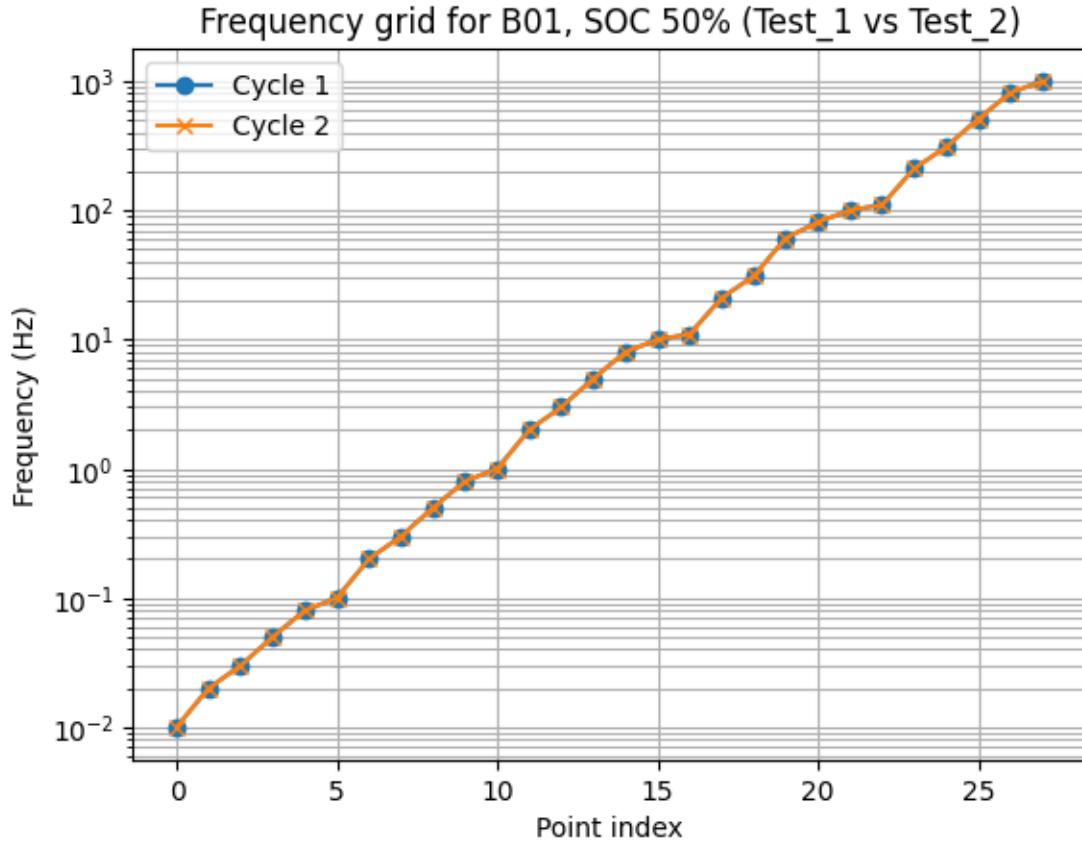
```
          cycle      soc
count    12320.0  12320.000
unique     NaN      NaN
top       NaN      NaN
freq      NaN      NaN
mean      1.5     52.500
std       0.5     28.833
min       1.0     5.000
25%       1.0     28.750
50%       1.5     52.500
75%       2.0     76.250
max       2.0    100.000
```

```
frequency_hz      0
z_real_ohm       0
z_imag_ohm       0
temperature_c     0
battery_id        0
cycle             0
soc               0
dtype: int64
```

```
np.int64(0)
```

Unique point counts per spectrum: [28]

```
count      440.0
mean      28.0
std       0.0
min      28.0
25%      28.0
50%      28.0
75%      28.0
max      28.0
Name: frequency_hz, dtype: float64
```



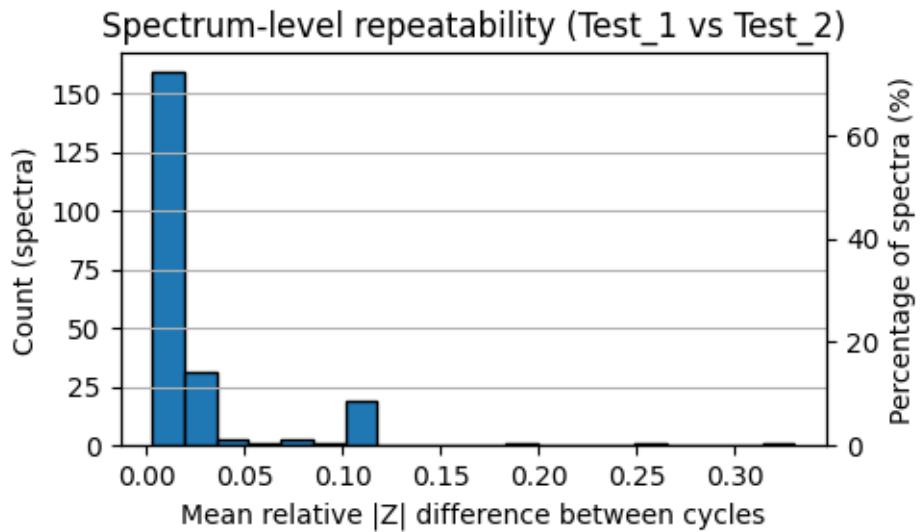
2. Cycle QC and averaging

Each cell and SOC has two EIS measurements (Test_1 and Test_2). Here I first check how consistent these repeats are, then average them to get one “denoised” spectrum per (battery_id, SOC) for the rest of the notebook.

```

count      220.000
mean       0.026
std        0.040
min        0.003
25%        0.009
50%        0.011
75%        0.025
max        0.331
Name: rel_diff, dtype: float64

```



The relative $|Z|$ difference between Test_1 and Test_2 is mostly very small (median 1.1%, 75% 2.5%), with a few larger outliers up to 33%.

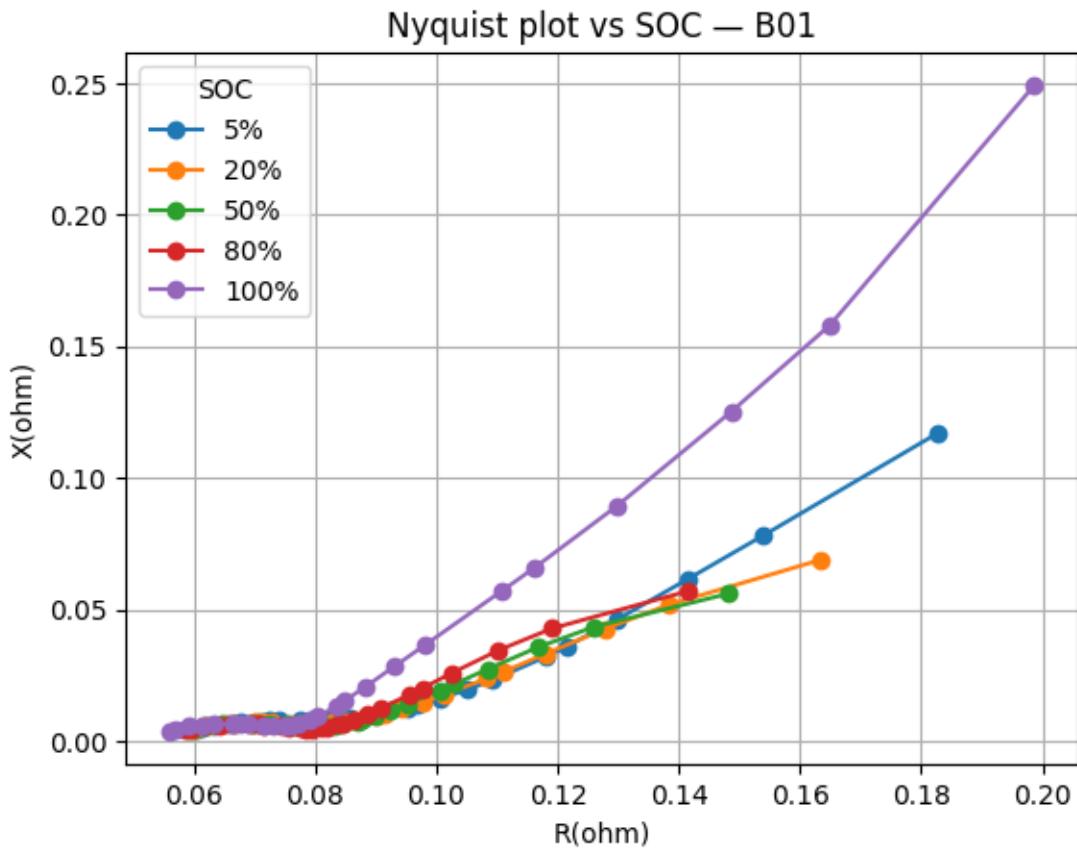
In this notebook I keep all spectra and simply average the two cycles per (battery_id, SOC), which is reasonable given the overall good repeatability. In a stricter setting one could drop spectra where the mean relative difference between cycles exceeds, for example, 20%.

Averaged EIS table shape: (6160, 6)

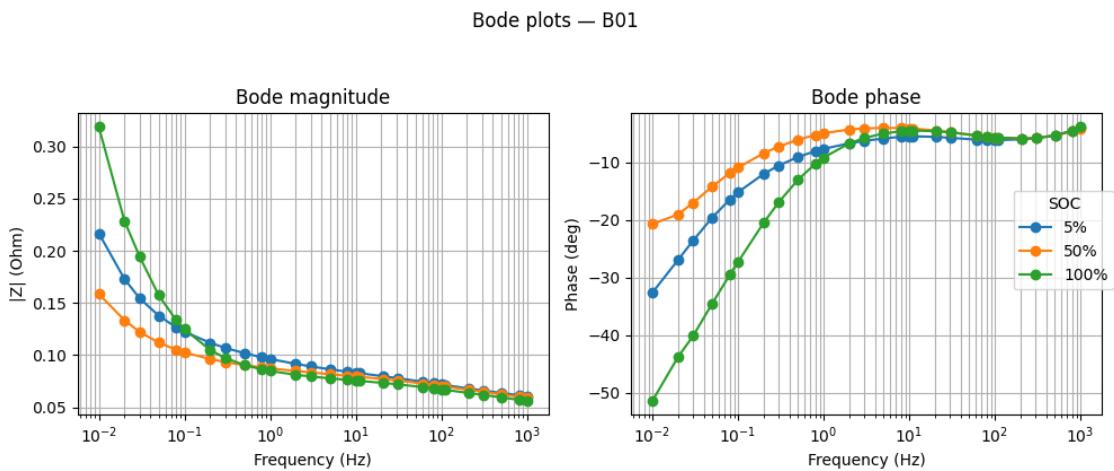
3. EIS visualisation vs SOC

3.1 Add magnitude and phase

3.2 Nyquist plots at different SOCs (one battery, one cycle)



3.3 Bode plots (magnitude & phase vs frequency)



4. Feature engineering: simple impedance fingerprints

Goal: compress each full EIS spectrum into a small feature vector, one row per (**battery_id**, **SOC**), that we can later feed to a simple ML model.

In this section I will:

- Create one feature row per (**battery_id**, **soc**).
- Use basic summary features (min/max/mean) of $|Z|$ and phase.
- Sample $|Z|$ and phase at a few fixed frequencies.

Why these engineered features?

The goal is to capture simple, physically meaningful aspects of the EIS curve without fitting a full equivalent circuit model:

- **zreal_lowfreq / zreal_highfreq**
 - Roughly “low-frequency / polarization resistance” vs “high-frequency / ohmic resistance”.
 - These shift with SOC as the cell’s effective resistance changes.
- **zmag_mean, zmag_max**
 - Overall “size” of the impedance arc and its peak.
 - Higher values usually mean higher resistance / lower conductivity.
- **phase_min, phase_max**
 - Range of capacitive / diffusive behaviour across frequency.
 - SOC changes how strongly capacitive or resistive the cell looks.
- **zmag_{f}Hz, phase_{f}Hz** at 0.01, 0.1, 1, 10, 100, 1000 Hz
 - These frequencies are chosen as simple, log-spaced points across the available 0.01–1000 Hz band.
 - Together they sample the low-, mid-, and high-frequency parts of the spectrum, giving a compact but informative “fingerprint” for SOC estimation.

4.1 Choose target frequencies from the measured grid I sample $|Z|$ and phase at a few log-spaced frequencies that exist in the data, to capture the overall shape of the spectrum in a compact way.

Measured frequency grid:

0: 0.01 Hz
1: 0.02 Hz
2: 0.03 Hz
3: 0.05 Hz
4: 0.08 Hz
5: 0.1 Hz
6: 0.2 Hz
7: 0.3 Hz
8: 0.5 Hz
9: 0.8 Hz
10: 1 Hz
11: 2 Hz

```

12: 3 Hz
13: 5 Hz
14: 8 Hz
15: 10 Hz
16: 11 Hz
17: 21 Hz
18: 31 Hz
19: 61 Hz
20: 81 Hz
21: 100 Hz
22: 110 Hz
23: 210 Hz
24: 310 Hz
25: 510 Hz
26: 810 Hz
27: 1000 Hz

```

Target frequencies:

```

0: 0.01 Hz
1: 0.1 Hz
2: 1 Hz
3: 10 Hz
4: 100 Hz
5: 1000 Hz

```

4.2 Build one feature vector per (battery_id, SOC) For each averaged EIS spectrum I create a single feature row.

I keep a few global summary statistics (extremes and averages), and I take $|Z|$ and phase directly at a set of fixed frequencies (0.01, 0.1, 1, 10, 100, 1000 Hz).

	battery_id	soc	zreal_lowfreq	zreal_highfreq	zmag_mean	zmag_max	\
0	B01	5	0.182653	0.060269	0.097467	0.216863	
1	B01	10	0.170783	0.059821	0.092930	0.187566	
2	B01	15	0.166635	0.059541	0.091769	0.180553	
3	B01	20	0.163275	0.059271	0.090735	0.177173	
4	B01	25	0.160309	0.059097	0.089930	0.172943	
	phase_min	phase_max	zmag_0p01Hz	phase_0p01Hz	zmag_0p1Hz	phase_0p1Hz	\
0	-32.621081	-4.004392	0.216863	-32.621081	0.122336	-15.181986	
1	-24.422087	-3.960382	0.187566	-24.422087	0.115161	-13.161971	
2	-22.644222	-3.972768	0.180553	-22.644222	0.113091	-12.632649	
3	-22.845436	-3.991457	0.177173	-22.845436	0.110929	-12.400498	
4	-22.036235	-4.016717	0.172943	-22.036235	0.109430	-12.100670	
	zmag_1p0Hz	phase_1p0Hz	zmag_10p0Hz	phase_10p0Hz	zmag_100p0Hz	\	
0	0.096208	-7.636961	0.083010	-5.468509	0.071766		
1	0.092719	-6.811176	0.081175	-4.923678	0.071196		
2	0.091900	-6.489224	0.080937	-4.856404	0.070999		

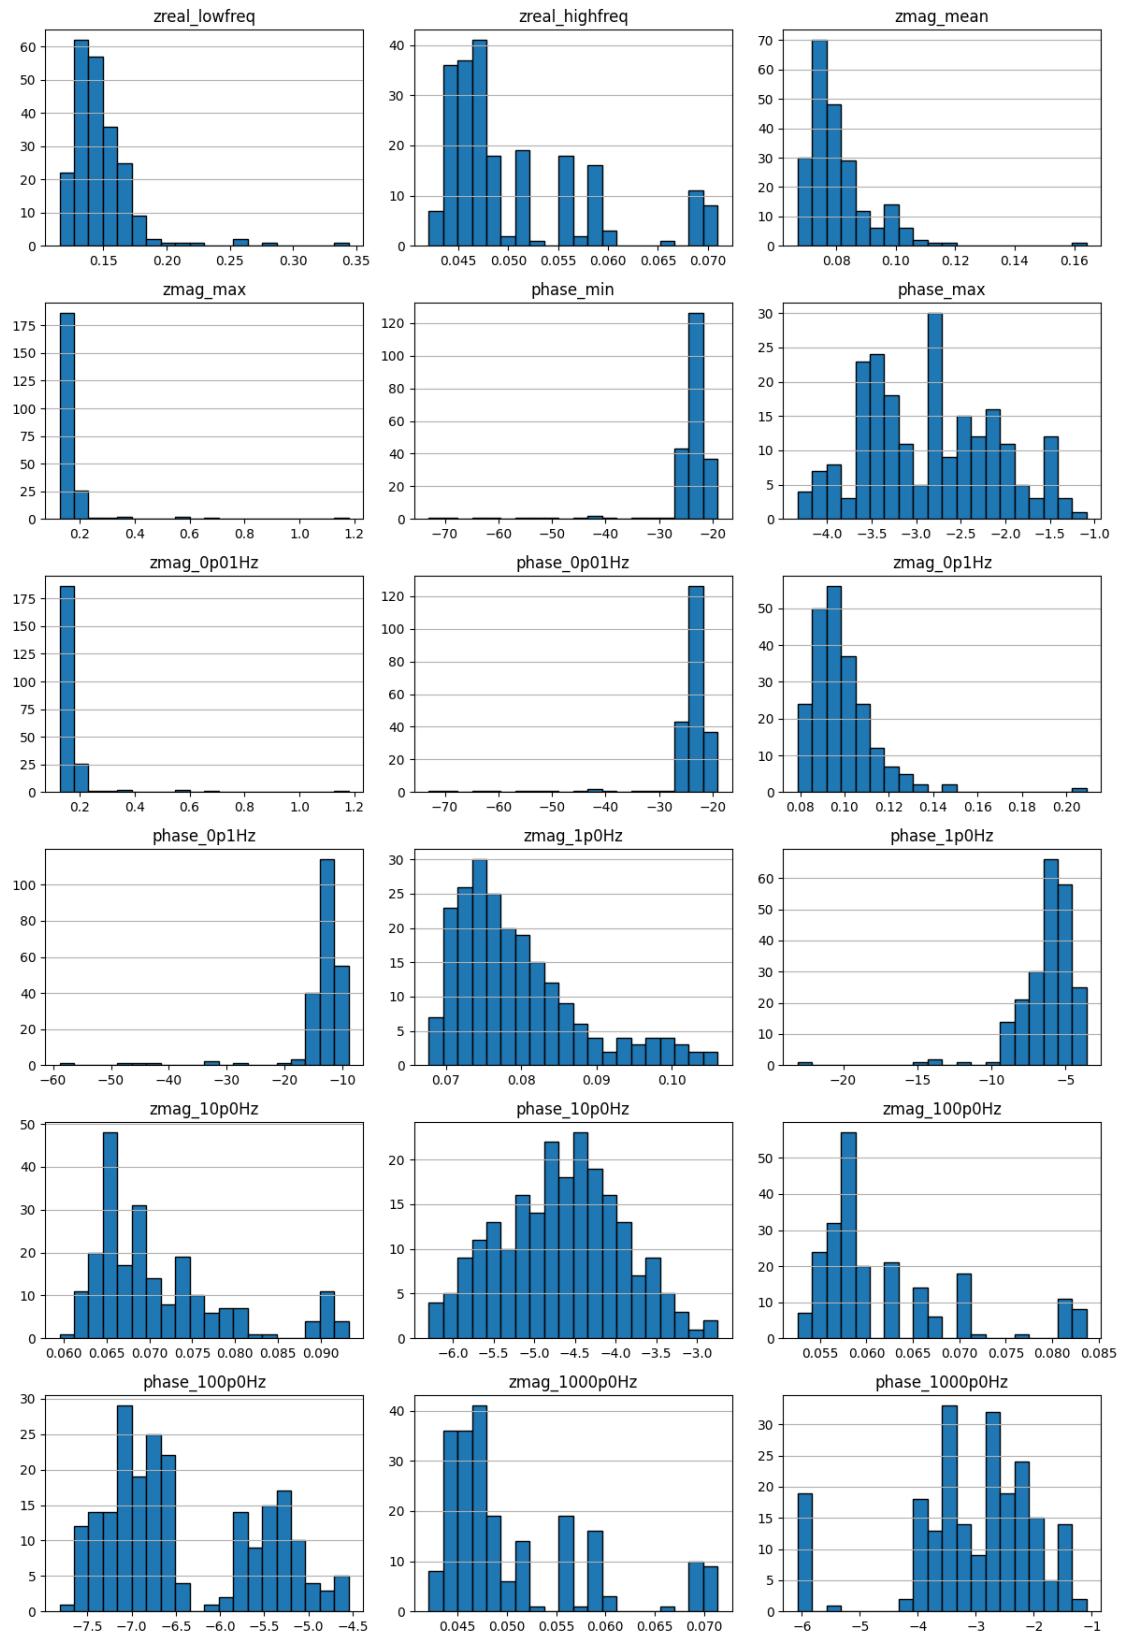
3	0.090892	-6.178033	0.080575	-4.715034	0.070769
4	0.090236	-5.955735	0.080347	-4.614360	0.070625

	phase_100p0Hz	zmag_1000p0Hz	phase_1000p0Hz
0	-6.121537	0.060417	-4.004392
1	-5.845588	0.059965	-3.960382
2	-5.846668	0.059685	-3.972768
3	-5.840951	0.059415	-3.991457
4	-5.832845	0.059243	-4.016717

4.3 Visualise key feature distributions Before modelling, I check a few key feature distributions to see:

- that values look reasonable,
- and that there is some spread across samples (so the model has signal to learn from).

Key feature distributions

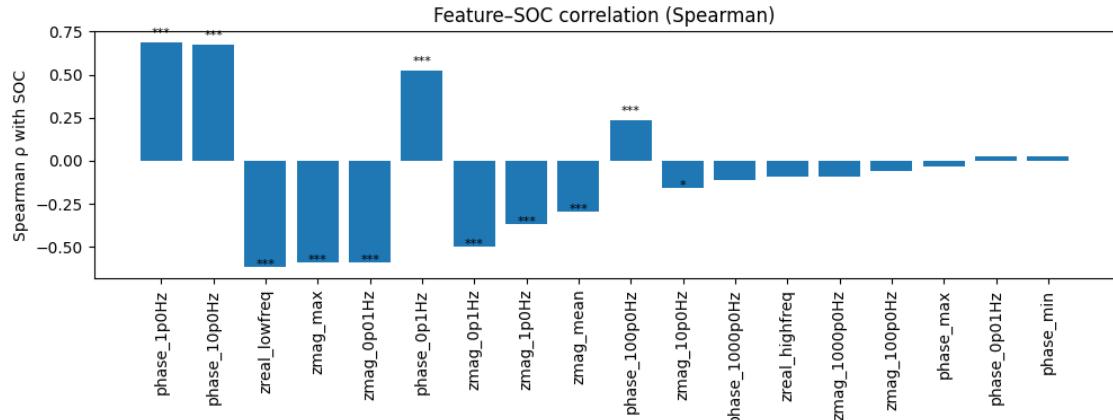


4.4 Correlation of engineered features with SOC To keep things readable, I only look at how each engineered feature correlates with **SOC** (rather than a full feature–feature heatmap).

I use **Spearman rank correlation** between SOC and all numeric features and mark statistically significant correlations:

- * : $p < 0.05$
- ** : $p < 0.01$
- ***: $p < 0.001$

This highlights which features are most informative for SOC and which ones are largely redundant.



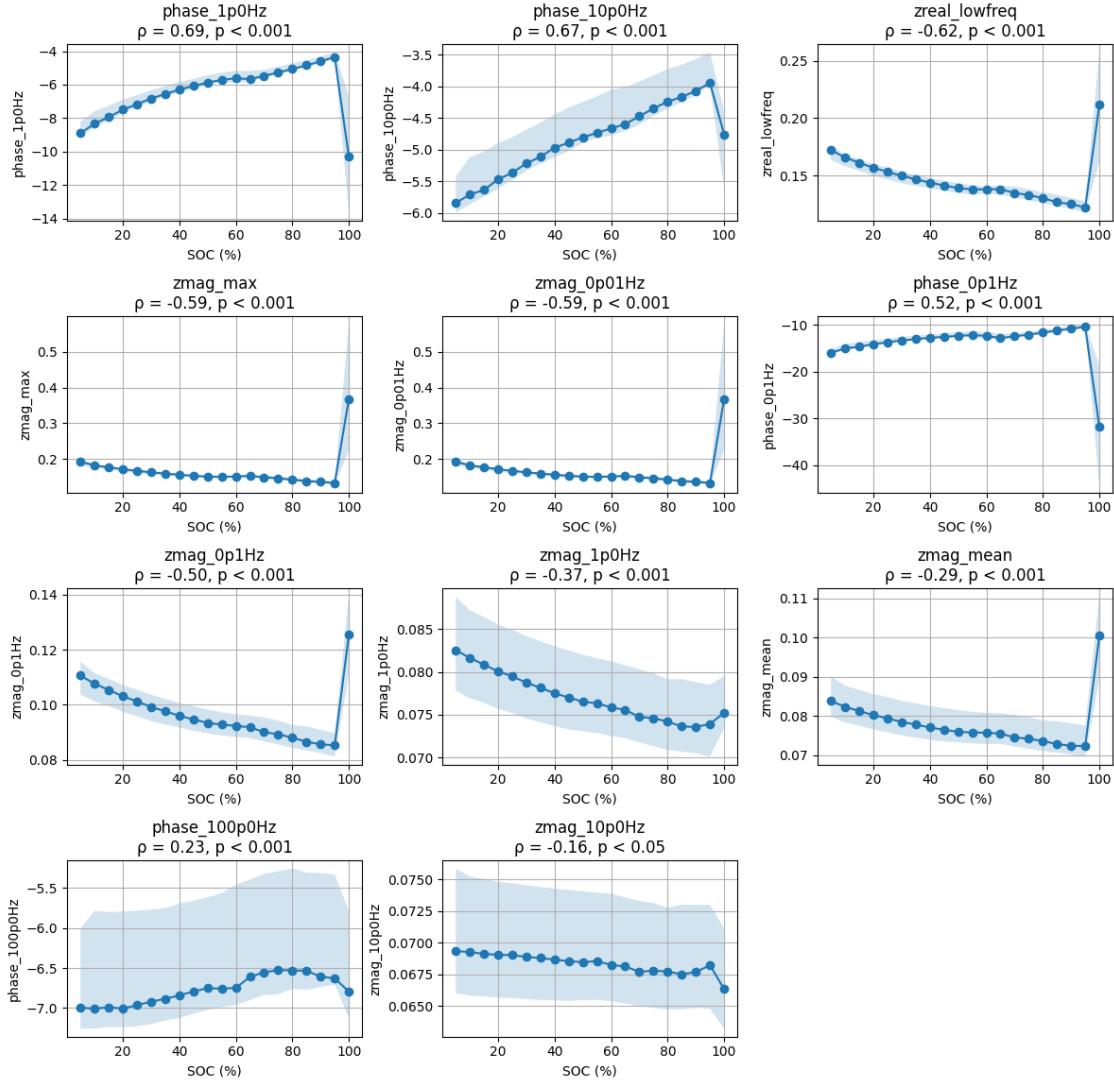
4.5 SOC vs significantly correlated features (aggregated by SOC) Scatter plots were noisy, so here I aggregate by SOC and plot, for each significant feature ($p < 0.05$):

- median value vs SOC (line),
- 25–75% range vs SOC (shaded band).

This makes monotonic trends much easier to see.

Significant features ($p < 0.05$): ['phase_1p0Hz', 'phase_10p0Hz', 'zreal_lowfreq', 'zmag_max', 'zmag_0p01Hz', 'phase_0p1Hz', 'zmag_0p1Hz', 'zmag_1p0Hz', 'zmag_mean', 'phase_100p0Hz', 'zmag_10p0Hz']

SOC vs significantly correlated features (aggregated by SOC)



5. SOC prediction target & problem setup

5.1 Define targets Main target (regression): - soc in %, e.g. 5–100.

Optional side target (bins): - Coarse SOC bands (0–20, 20–40, 40–60, 60–80, 80–100%) for an easier “are we in the right band?” view.

```
soc_bin
0    44
1    44
2    44
3    44
```

4 44

Name: count, dtype: int64

5.2 Feature subsets: all vs significant

I compare two feature sets:

- **all**: all engineered numeric features except identifiers and targets
- **sig**: only features with significant Spearman correlation to SOC ($p < 0.05$)

Both will be evaluated inside the same nested GroupKFold CV.

All features: 18

Significant features: 11

Significant feature names: ['zreal_lowfreq', 'zmag_mean', 'zmag_max',
'zmag_0p01Hz', 'zmag_0p1Hz', 'phase_0p1Hz', 'zmag_1p0Hz', 'phase_1p0Hz',
'zmag_10p0Hz', 'phase_10p0Hz', 'phase_100p0Hz']

5.3 Models and nested GroupKFold setup

I compare two models:

- **RandomForestRegressor** (non-linear, tree ensemble)
- **ElasticNet** (linear model with L1/L2 regularisation in a scaled pipeline)

Nested grouped CV:

- **Outer**: GroupKFold($n_splits=5$) on **battery_id**
 - ~80% of cells for training, ~20% for testing in each fold.
- **Inner**: GroupKFold($n_splits=3$) on the outer-train cells
 - used only for hyperparameter tuning (MAE as objective).

5.4 Nested GroupKFold CV: run and collect metrics

For each outer fold (5 folds):

- Use only the outer-train cells in a 3-fold GroupKFold inner CV to tune hyperparameters for each (feature set, model) combination.
- Refit the best model on all outer-train data.
- Evaluate on the outer-test cells:
 - MAE (mean absolute error) in SOC %
 - RMSE (root mean squared error) in SOC %
 - SOC-bin accuracy (using the coarse bins defined earlier)

I also record how many samples and batteries are in train/test per fold.

Nested CV: 0% | 0/20 [00:00<?, ?it/s]

```
Fold 0, fs=all, model=rf | MAE=3.27, RMSE=4.88, bin_acc=81.67%,  
best={'max_depth': None, 'n_estimators': 300}  
Fold 0, fs=all, model=enet | MAE=4.02, RMSE=6.90, bin_acc=81.67%,  
best={'model__alpha': 0.1, 'model__l1_ratio': 1.0}  
Fold 0, fs=sig, model=rf | MAE=3.86, RMSE=5.11, bin_acc=83.33%,  
best={'max_depth': None, 'n_estimators': 300}  
Fold 0, fs=sig, model=enet | MAE=4.79, RMSE=8.13, bin_acc=80.00%,  
best={'model__alpha': 0.1, 'model__l1_ratio': 1.0}
```

```

Fold 1, fs=all, model=rf | MAE=3.75, RMSE=5.39, bin_acc=90.00%,
best={'max_depth': None, 'n_estimators': 150}
Fold 1, fs=all, model=enet | MAE=4.05, RMSE=5.98, bin_acc=87.50%,
best={'model__alpha': 0.1, 'model__l1_ratio': 1.0}
Fold 1, fs=sig, model=rf | MAE=4.41, RMSE=6.56, bin_acc=85.00%,
best={'max_depth': None, 'n_estimators': 150}
Fold 1, fs=sig, model=enet | MAE=4.11, RMSE=5.43, bin_acc=85.00%,
best={'model__alpha': 0.1, 'model__l1_ratio': 1.0}
Fold 2, fs=all, model=rf | MAE=4.45, RMSE=7.33, bin_acc=72.50%,
best={'max_depth': None, 'n_estimators': 300}
Fold 2, fs=all, model=enet | MAE=2.89, RMSE=3.74, bin_acc=85.00%,
best={'model__alpha': 0.1, 'model__l1_ratio': 1.0}
Fold 2, fs=sig, model=rf | MAE=4.53, RMSE=10.12, bin_acc=70.00%,
best={'max_depth': None, 'n_estimators': 300}
Fold 2, fs=sig, model=enet | MAE=2.83, RMSE=3.96, bin_acc=77.50%,
best={'model__alpha': 0.1, 'model__l1_ratio': 1.0}
Fold 3, fs=all, model=rf | MAE=10.59, RMSE=14.73, bin_acc=55.00%,
best={'max_depth': None, 'n_estimators': 300}
Fold 3, fs=all, model=enet | MAE=2.62, RMSE=3.32, bin_acc=80.00%,
best={'model__alpha': 0.1, 'model__l1_ratio': 1.0}
Fold 3, fs=sig, model=rf | MAE=12.89, RMSE=18.38, bin_acc=52.50%,
best={'max_depth': None, 'n_estimators': 300}
Fold 3, fs=sig, model=enet | MAE=3.66, RMSE=4.71, bin_acc=82.50%,
best={'model__alpha': 0.1, 'model__l1_ratio': 1.0}
Fold 4, fs=all, model=rf | MAE=7.56, RMSE=11.39, bin_acc=60.00%,
best={'max_depth': None, 'n_estimators': 150}
Fold 4, fs=all, model=enet | MAE=5.89, RMSE=8.62, bin_acc=67.50%,
best={'model__alpha': 0.1, 'model__l1_ratio': 1.0}
Fold 4, fs=sig, model=rf | MAE=7.69, RMSE=11.97, bin_acc=60.00%,
best={'max_depth': None, 'n_estimators': 300}
Fold 4, fs=sig, model=enet | MAE=7.80, RMSE=16.60, bin_acc=65.00%,
best={'model__alpha': 0.1, 'model__l1_ratio': 1.0}

```

	fold	n_train_samples	n_test_samples	n_train_batteries	n_test_batteries	\
0	0	160	60	8	3	
1	1	180	40	9	2	
2	2	180	40	9	2	
3	3	180	40	9	2	
4	4	180	40	9	2	
				train_batteries	test_batteries	
0		B02, B03, B04, B05, B07, B08, B09, B10		B01, B06, B11		
1		B01, B02, B03, B04, B06, B07, B08, B09, B11			B05, B10	
2		B01, B02, B03, B05, B06, B07, B08, B10, B11			B04, B09	
3		B01, B02, B04, B05, B06, B07, B09, B10, B11			B03, B08	
4		B01, B03, B04, B05, B06, B08, B09, B10, B11			B02, B07	

5.5 Compact summary of performance I summarise performance across the 5 outer folds by feature set and model:

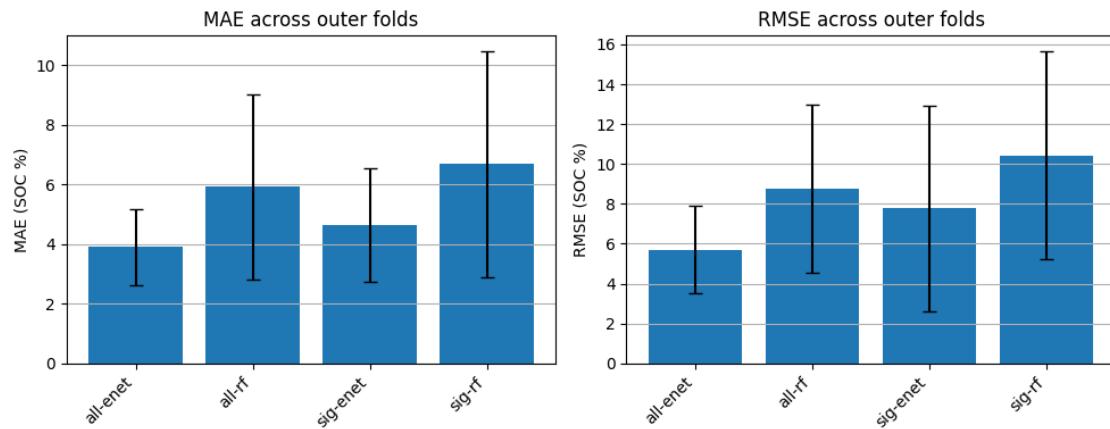
- Mean \pm std of MAE and RMSE in SOC %
- Mean SOC-bin accuracy

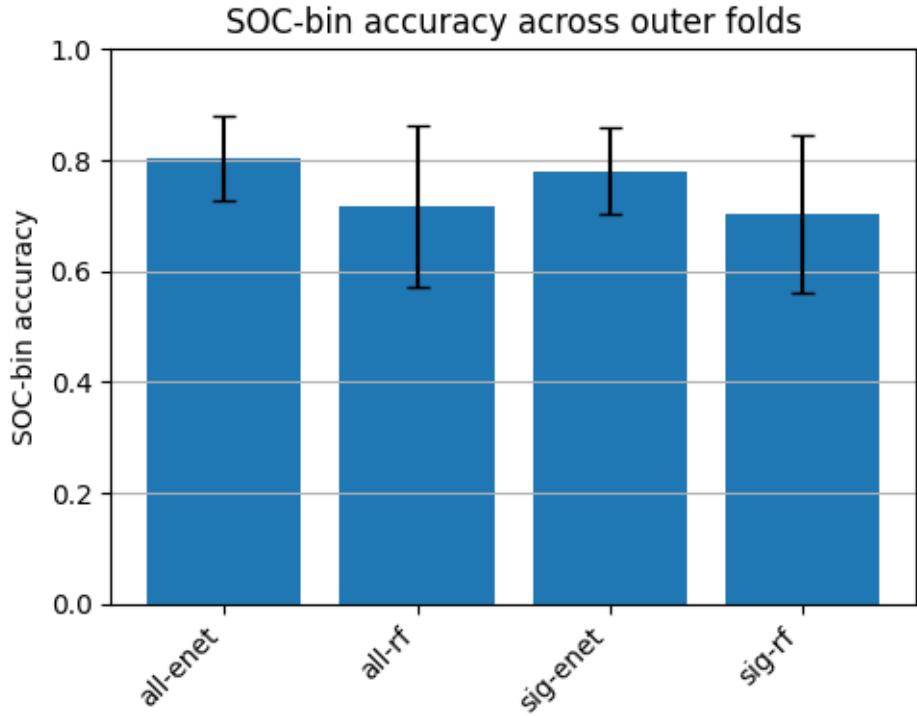
This lets me compare:

- RandomForest vs ElasticNet (non-linear vs linear)
- All features vs only the statistically significant ones

Summary metrics across outer folds:

	feature_set	model	mae_mean	mae_std	rmse_mean	rmse_std	bin_acc_mean	bin_acc_std
0	all	enet	3.897597	1.289881	5.711748	2.211639	0.803333	0.077415
1	all	rf	5.923722	3.097528	8.742880	4.215280	0.718333	0.145821
2	sig	enet	4.637253	1.906821	7.763793	5.182472	0.780000	0.077862
3	sig	rf	6.676111	3.784897	10.429521	5.219661	0.701667	0.142205





6. Summary of SOC modelling experiment

This section briefly sums up the SOC prediction experiment based on the engineered EIS features and the nested, grouped cross-validation.

6.1 Aggregated metrics table First I summarise the nested GroupKFold results across outer folds for each (feature set, model) combination.

```
<pandas.io.formats.style.Styler at 0x7265add66e50>
```

6.2 Setup (reminder)

- **Input:** engineered EIS features per (battery_id, SOC), including:
 - global stats: zreal_lowfreq, zreal_highfreq, zmag_mean, zmag_max, phase_min, phase_max
 - |Z| and phase at 0.01, 0.1, 1, 10, 100, 1000 Hz
- **Target:** SOC in %, evaluated as:
 - **regression:** MAE / RMSE in SOC %
 - **coarse SOC bins:** [0–20, 20–40, 40–60, 60–80, 80–100]% for bin accuracy
- **Splits:** nested, grouped by battery_id
 - **Outer:** GroupKFold(n_splits=5)
 - ~80% of cells for training, ~20% for testing in each fold

- **Inner:** GroupKFold(n_splits=3) on outer-train cells
→ used only for hyperparameter tuning (MAE as objective)
- **Feature sets:**
 - **all** – all engineered features (except identifiers and targets)
 - **sig** – only features with Spearman p < 0.05 vs SOC
- **Models:**
 - **rf** – RandomForestRegressor
 - **enet** – ElasticNet (in a scaled pipeline, tuned `alpha` and `l1_ratio`)

The `summary` table aggregates metrics across the 5 outer folds for each (feature set, model) combination.

6.3 Quantitative results (mean ± std across outer folds)

From the aggregated results (`summary`):

- **all + enet** (ElasticNet on all features):
 - MAE **3.9 ± 1.3 % SOC**
 - RMSE **5.7 ± 2.2 % SOC**
 - SOC-bin accuracy **0.80 ± 0.08**
- **all + rf** (RandomForest on all features):
 - MAE **5.9 ± 3.1 % SOC**
 - RMSE **8.7 ± 4.2 % SOC**
 - SOC-bin accuracy **0.72 ± 0.15**
- **sig + enet** (ElasticNet on significant features only):
 - MAE **4.6 ± 1.9 % SOC**
 - RMSE **7.8 ± 5.2 % SOC**
 - SOC-bin accuracy **0.78 ± 0.08**
- **sig + rf** (RandomForest on significant features only):
 - MAE **6.7 ± 3.8 % SOC**
 - RMSE **10.4 ± 5.2 % SOC**
 - SOC-bin accuracy **0.70 ± 0.14**

6.4 Main observations

1. Best overall configuration

- The best-performing configuration is **ElasticNet with all features** (`feature_set = "all", model = "enet"`):
 - MAE **3.9 % SOC**
 - RMSE **5.7 % SOC**
 - SOC-bin accuracy **80%** (correct 20%-wide SOC band)

- This is the natural “baseline SOC model” to highlight in this notebook.

2. Linear vs non-linear model

- On both feature sets, **ElasticNet beats RandomForest**:
 - With all features:
 - * MAE improves from $\sim 5.9 \rightarrow \sim 3.9\%$ SOC
 - * RMSE improves from $\sim 8.7 \rightarrow \sim 5.7\%$ SOC
 - * SOC-bin accuracy improves from $\sim 72\% \rightarrow \sim 80\%$
 - Interpretation:
 - The EIS features are smooth and highly correlated.
 - A regularised **linear model** already captures most of the SOC signal; the non-linear RandomForest does not add value here and may overfit given only 11 cells.

3. All features vs “significant-only” features

- For ElasticNet:
 - Restricting to only the “significant” features (Spearman $p < 0.05$) **slightly worsens** MAE, RMSE and bin accuracy versus using all features.
- For RandomForest:
 - The same pattern holds: using only the significant subset is not better.
- Interpretation:
 - Dropping features based purely on per-feature correlation with SOC does not help in this case.
 - Given strong multicollinearity and the use of regularisation, it is actually safe (and slightly better) to keep the full engineered feature set.

4. Variability across folds

- For the best config (all + enet), variability across outer folds is moderate:
 - MAE std **1.3 % SOC**
 - RMSE std **2.2 % SOC**
 - SOC-bin accuracy std **0.08**
- This reflects the small number of cells (11) and the fact that each fold holds out a different subset of batteries.
- Despite this, the ranking of configurations is consistent: **all+enet** best, then **sig+enet**, both RandomForest variants clearly worse.

6.5 Takeaways

Overall, the experiment shows that:

- With a compact set of physically motivated EIS features, a **regularised linear model (ElasticNet)** can estimate SOC with:
 - typical errors around **4% SOC** (MAE),
 - and **$\sim 80\%$ accuracy** in coarse 20%-wide SOC bands,
 - when evaluated with 5-fold **GroupKFold** that holds out entire cells.
- A more complex **RandomForest** model does not improve performance on this dataset and is consistently worse across all metrics.
- Restricting to only “statistically significant” features (Spearman $p < 0.05$) slightly hurts performance; the full engineered feature set works best, likely because:

- redundant but informative features are handled well by ElasticNet’s regularisation,
- and single-feature correlation is a crude way to do feature selection on smooth spectral data.

Given the small number of cells (11), these results are a **proof-of-concept** rather than a final production model, but they already demonstrate that:

Impedance fingerprints contain enough information to recover SOC reasonably well, and simple, interpretable models can leverage that information effectively.