

# 03 – Healthy vs Aged Classification from EIS

Chouaib (github.com/ChouaibB)

December 19, 2025

## Contents

1.	Overview & data choice for healthy vs aged classification . . . . .	2
2.	Dataset and feature pipeline (from 02 – SoH from EIS) . . . . .	2
2.1	Load EIS data . . . . .	3
2.2	Build the impedance feature table for ML . . . . .	3
3.	SoH-based labels for healthy vs aged classification . . . . .	4
3.1	Define SoH-based binary labels . . . . .	4
3.2	Inspect healthy vs aged class balance . . . . .	4
4.	Modelling: nested CV with logistic regression, SVM and GP . . . . .	5
4.1	Features, target and groups . . . . .	5
4.2	Define models and simple hyperparameter grids . . . . .	6
4.3	Main models training loop . . . . .	6
5.	Results: outer-fold performance . . . . .	7
5.1	Boxplots of ROC AUC and PR AUC . . . . .	8
5.2	ROC and precision–recall curves . . . . .	8
5.3	Interpretation of healthy vs aged classification results . . . . .	8
6.	Feature importance and 2D decision regions . . . . .	9
6.1	Final models and permutation-based feature importance . . . . .	10
6.2	Decision regions in a shared 2D feature plane . . . . .	10

## 1. Overview & data choice for healthy vs aged classification

For this notebook I originally planned to use the **NASA PCoE Li-ion Battery Aging Datasets** (<https://data.nasa.gov/dataset/li-ion-battery-aging-datasets>), which contain full aging trajectories down to clear **end-of-life (EoL)** and are widely used for **SoH / RUL** studies. At the time of this work I could not access that dataset, so I instead reuse the SoH-from-EIS dataset from **02 – SoH from EIS (Rashid et al.)** as a **stand-in** to demonstrate a simple healthy vs aged classification workflow.

In real applications, a cell is typically considered to have reached EoL when its capacity (or SoH) drops to around **80%** of nominal. The Rashid SoH-from-EIS dataset I use here only covers SoH levels from **100% down to 80%** in 5% steps, and does **not** include cells degraded below 80% like the NASA PCoE data would. To still illustrate a binary screening task with the data at hand, I define:

- **healthy** = SoH 90%
- **aged** = SoH < 90%

This threshold is therefore **chosen for demonstration**, not as a universal definition of EoL, and should be interpreted as “early screening for cells that have started to age” rather than a hard safety or warranty limit. In practice, I would expect cells aged **below 80% SoH to develop a noticeably different impedance fingerprint/signature** from the SoH ranges considered here (SoH 90% and SoH < 90%), which could **substantially change** both the models and the conclusions. The current notebook should therefore be read as a **workflow demonstration** under limited data, not as a definitive analysis of fully aged cells.

---

## 2. Dataset and feature pipeline (from 02 – SoH from EIS)

As in

**02 – SoH from EIS (Rashid et al.)**, I use the public dataset:

Rashid, Muhammad; Faraji-Niri, Mona; Sansom, Jonathan; Sheikh, Muhammad; Widanage, Dhammadika; Marco, James (2023),  
**“Dataset for rapid state of health estimation of lithium batteries using EIS and machine learning: Training and validation”**,  
*Data in Brief*, 48, 109157, doi: 10.1016/j.dib.2023.109157.  
Available at: <https://data.mendeley.com/datasets/mn9fb7wdx6/3>  
Original data: **“DIB\_Data”**, Mendeley Data, V3, doi: 10.17632/mn9fb7wdx6.3 (CC0 1.0).

Dataset highlights (same as in 02 – SoH from EIS):

- **25 cylindrical Li-ion cells**, aged from SoH 100% down to 80% in 5% steps (100, 95, 90, 85, 80%).
- At each SoH stage, reference performance tests (capacity / SoH) plus **electrochemical impedance spectroscopy (EIS)** at multiple **SOC** and **temperature** conditions.
- Designed specifically to study **fast SoH estimation from EIS with machine learning**.

In this notebook I **reuse the same data loading, cleaning and feature engineering pipeline**

as in 02 – SoH from EIS:

- I load the raw EIS spectra and associated metadata,
- perform the same basic QC and filtering,
- and construct the same engineered impedance feature table (ohmic / low-frequency resistances, summary  $|Z|$  and phase statistics, and sampled spectral points).

For details of the data preparation steps, please refer to

**02 – SoH from EIS (Rashid et al.).** Here I start directly from the prepared feature table and focus on the **binary classification** of healthy vs aged cells based on those impedance features.

## 2.1 Load EIS data

```
n_cells  n_eis_tests          soh_levels          soc_levels temp_levels \
0        24            360  80, 85, 90, 95, 100  5, 20, 50, 70, 95  15, 25, 35

freq_points_per_spectrum
0                      61
```

## 2.2 Build the impedance feature table for ML

Feature table shape: (360, 27)

```
cell_id  soh_pct  temp_c  soc_pct  R_hf_ohm  R_lf_ohm  delta_R_ohm \
0        2         95      15       5   0.02995   0.07366   0.04371
1        2         95      15       20  0.03001   0.07916   0.04915
2        2         95      15       50  0.02965   0.03573   0.00608
3        2         95      15       70  0.02960   0.03803   0.00843
4        2         95      15       95  0.02958   0.04171   0.01213

Zmag_mean  Zmag_std  Zmag_min ...  Zmag_f1p0  phase_f1p0  Zmag_f10p0 \
0  0.040826  0.016416  0.024201 ...  0.042847  -0.282403  0.031430
1  0.041919  0.018060  0.024285 ...  0.042731  -0.290515  0.031688
2  0.029907  0.003850  0.023819 ...  0.030956  -0.031686  0.029673
3  0.030187  0.004274  0.023747 ...  0.030942  -0.037886  0.029639
4  0.031920  0.005565  0.023771 ...  0.034760  -0.067860  0.030048

phase_f10p0  Zmag_f100p0  phase_f100p0  Zmag_f1000p0  phase_f1000p0 \
0  -0.123542  0.027523  -0.077613  0.024205  0.097449
1  -0.122719  0.027668  -0.079707  0.024291  0.095818
2  -0.047501  0.026807  -0.059273  0.023819  0.103972
3  -0.051610  0.026685  -0.058380  0.023747  0.103440
4  -0.092483  0.026772  -0.062605  0.023772  0.101259

Zmag_f10000p0  phase_f10000p0
0    0.043538  0.812199
1    0.043579  0.811200
2    0.043398  0.818643
3    0.043320  0.818540
```

```
4          0.043262        0.817930
```

```
[5 rows x 27 columns]
```

### 3. SoH-based labels for healthy vs aged classification

In this section I convert the SoH information into a **binary health label** that will be used as the target for the classifier:

- **healthy** = SoH 90%
- **aged** = SoH < 90%


I first add a continuous SoH fraction (`soh_frac`) and then derive the binary `health_label` (and a numeric `target_aged` = 1 for aged, 0 for healthy). I also inspect the resulting **class balance** (counts and percentages) to see how many samples fall into each group before training any models.

#### 3.1 Define SoH-based binary labels

```
  cell_id  soh_pct  temp_c  soc_pct  health_label
0         2       95      15        5    healthy
1         2       95      15       20    healthy
2         2       95      15       50    healthy
3         2       95      15       70    healthy
4         2       95      15       95    healthy
```

#### 3.2 Inspect healthy vs aged class balance

```
  health_label  n_samples   pct
0     healthy        210  58.3
1       aged         150  41.7
```



## 4. Modelling: nested CV with logistic regression, SVM and GP

In this section I train three **probabilistic classifiers** on the original engineered impedance features:

1. Regularised logistic regression (`LogisticRegression`)
2. RBF-kernel SVM (`SVC(kernel="rbf", probability=True)`)
3. Gaussian Process Classifier (`GaussianProcessClassifier`)

I use a **nested cross-validation** scheme:

- **Outer loop:** 5-fold grouped splits by `cell_id` to estimate performance variability across different groups of cells.
- **Inner loop:** grouped K-fold CV (here 3 folds) on the outer-training data for hyperparameter tuning via `GridSearchCV`.

All models are wrapped in **Pipelines** (with `StandardScaler` where needed) and evaluated using **ROC AUC** and **PR AUC** on the outer test folds. The results are collected in a single table for later summary and visualisation.

### 4.1 Features, target and groups

`((360, 25), (360,))`

## 4.2 Define models and simple hyperparameter grids

## 4.3 Main models training loop

### 4.3a Define CV objects and inspect outer-fold class balance

outer_fold	train_n	train_neg	train_pos	test_n	test_neg	test_pos
0	1	285	180	105	75	30
1	2	285	165	120	75	45
2	3	285	165	120	75	45
3	4	285	150	135	75	60
4	5	300	180	120	60	30

### 4.3b Model training loop (using the same CV objects)

==== Model: log\_reg ===

```
log_reg | outer CV: 0% | 0/5 [00:00<?, ?it/s]

Outer fold 1/5
    best inner ROC AUC = 0.963 | outer ROC AUC = 0.834, PR AUC = 0.905
Outer fold 2/5
    best inner ROC AUC = 0.952 | outer ROC AUC = 1.000, PR AUC = 1.000
Outer fold 3/5
    best inner ROC AUC = 0.962 | outer ROC AUC = 0.996, PR AUC = 0.994
Outer fold 4/5
    best inner ROC AUC = 0.978 | outer ROC AUC = 0.903, PR AUC = 0.687
Outer fold 5/5
    best inner ROC AUC = 0.935 | outer ROC AUC = 1.000, PR AUC = 1.000
```

==== Model: svm\_rbf ===

```
svm_rbf | outer CV: 0% | 0/5 [00:00<?, ?it/s]

Outer fold 1/5
    best inner ROC AUC = 0.964 | outer ROC AUC = 0.689, PR AUC = 0.751
Outer fold 2/5
    best inner ROC AUC = 0.909 | outer ROC AUC = 1.000, PR AUC = 1.000
Outer fold 3/5
    best inner ROC AUC = 0.977 | outer ROC AUC = 0.604, PR AUC = 0.695
Outer fold 4/5
    best inner ROC AUC = 0.996 | outer ROC AUC = 0.899, PR AUC = 0.638
Outer fold 5/5
    best inner ROC AUC = 0.938 | outer ROC AUC = 1.000, PR AUC = 1.000
```

==== Model: gp\_classifier ===

```
gp_classifier | outer CV: 0% | 0/5 [00:00<?, ?it/s]

Outer fold 1/5
    best inner ROC AUC = 0.962 | outer ROC AUC = 0.861, PR AUC = 0.908
Outer fold 2/5
```

```

best inner ROC AUC = 0.941 | outer ROC AUC = 0.997, PR AUC = 0.995
Outer fold 3/5
best inner ROC AUC = 0.963 | outer ROC AUC = 0.649, PR AUC = 0.712
Outer fold 4/5
best inner ROC AUC = 0.961 | outer ROC AUC = 0.927, PR AUC = 0.730
Outer fold 5/5
best inner ROC AUC = 0.933 | outer ROC AUC = 0.946, PR AUC = 0.940

```

	outer_fold	model	roc_auc	pr_auc	best_inner_roc_auc	\
0	2	log_reg	1.000000	1.000000	0.952222	
1	5	log_reg	1.000000	1.000000	0.935370	
2	2	svm_rbf	1.000000	1.000000	0.908593	
3	5	svm_rbf	1.000000	1.000000	0.938210	
4	2	gp_classifier	0.997037	0.995474	0.940704	
5	3	log_reg	0.995556	0.993845	0.962296	
6	5	gp_classifier	0.945556	0.939726	0.933457	
7	4	gp_classifier	0.926667	0.729529	0.961420	
8	4	log_reg	0.903333	0.687388	0.978333	
9	4	svm_rbf	0.899444	0.638067	0.995556	
10	1	gp_classifier	0.860741	0.908004	0.961811	
11	1	log_reg	0.834074	0.905359	0.963457	
12	1	svm_rbf	0.688889	0.750990	0.964115	
13	3	gp_classifier	0.648889	0.712016	0.962741	
14	3	svm_rbf	0.604444	0.694608	0.977481	

	best_params	n_test_samples
0	{'clf__C': 0.1}	75
1	{'clf__C': 0.1}	60
2	{'clf__C': 1.0, 'clf__gamma': 0.01}	75
3	{'clf__C': 1.0, 'clf__gamma': 0.01}	60
4	{'clf__kernel': RBF(length_scale=10)}	75
5	{'clf__C': 0.1}	75
6	{'clf__kernel': RBF(length_scale=1)}	60
7	{'clf__kernel': RBF(length_scale=10)}	75
8	{'clf__C': 1.0}	75
9	{'clf__C': 10.0, 'clf__gamma': 0.01}	75
10	{'clf__kernel': RBF(length_scale=10)}	75
11	{'clf__C': 100.0}	75
12	{'clf__C': 1.0, 'clf__gamma': 'scale'}	75
13	{'clf__kernel': RBF(length_scale=1)}	75
14	{'clf__C': 0.1, 'clf__gamma': 1.0}	75

## 5. Results: outer-fold performance


First I summarise the nested CV performance across the 5 outer folds for each model using boxplots of:

- ROC AUC (aged vs healthy)

- PR AUC (precision–recall for the aged class)


Each point comes from one outer test fold; the boxes show the spread across folds for each model.

## 5.1 Boxplots of ROC AUC and PR AUC



## 5.2 ROC and precision–recall curves

To visualise the trade-offs, I build ROC and PR curves for each model by concatenating the predictions from all outer test folds and computing the curves on these cross-validated predictions.



## 5.3 Interpretation of healthy vs aged classification results

### Class balance & splits

- Each outer test fold has **60–75 samples** with both classes present (e.g. test sets range from 30/45 to 60/15 healthy/aged), so ROC and PR metrics are well-defined and there is no obvious class-collapse in any fold.

## Overall separability

- All three models achieve **high ROC AUC (~0.84–0.95 on average)** and **PR AUC (~0.82–0.92)** across the 5 outer folds, which means the impedance features carry a strong signal for the **healthy vs aged** label.
- The pooled ROC/PR curves show that, for a wide range of thresholds, we can keep the **false positive rate reasonably low** while maintaining good recall for aged cells.

## Model comparison

- **Logistic regression**
  - Highest and most stable performance: mean ROC AUC **0.95**, PR AUC **0.92** with relatively small fold-to-fold spread.
  - Several outer folds reach ROC/PR **1.0**, which suggests that in those particular cell splits the classes are almost perfectly linearly separable in the chosen feature space.
  - Given its simplicity, this is a very strong baseline and indicates that a mostly linear decision boundary in impedance features is already very effective for this healthy/aged split.
- **Gaussian Process Classifier**
  - Close runner-up: mean ROC AUC **0.88**, PR AUC **0.86**.
  - Shows more variability across folds (one fold drops to ROC AUC **0.65**).
- **RBF SVM**
  - Slightly weaker and more variable: mean ROC AUC **0.84**, PR AUC **0.82**, with one difficult fold around ROC AUC **0.60** and others close to 1.0.
  - The higher variance across folds likely reflects sensitivity to the hyperparameters and to which cells land in train vs test.

## Caveats

- The near-perfect scores on some folds are plausible given:
  - relatively small outer test sets (60–75 samples), and
  - a **constructed label** (SoH **90%** vs **< 90%**) that may align quite cleanly with certain impedance patterns for some cells.
- However, the fact that:
  - the evaluation is **nested**,
  - splits are grouped by **cell\_id**, and
  - not all folds achieve 1.0, suggests we are not seeing a trivial leak, but rather a genuinely strong separability in this dataset and threshold choice.

Overall, the results show that impedance-based features in this Rashid et al. dataset are highly informative for the early **healthy vs aged** screening task, with **regularised logistic regression** already providing a very competitive and stable classifier, and SVM / GP offering alternative non-linear and probabilistic views with slightly lower but comparable performance.

## 6. Feature importance and 2D decision regions

Here I look at **which impedance features matter most** for the healthy vs aged classification, and how the three classifiers behave in a **2D feature plane**.


- First, I refit each model once on the **full dataset** (using grouped CV on `cell_id` for hyper-parameter tuning) and compute **permutation importance**, using ROC AUC as the scoring metric.
- Then, for each model, I take its **two most important features**, train a small 2D version of the classifier on those features, and visualise the **decision regions / probability maps** together with the data points.

## 6.1 Final models and permutation-based feature importance

```
==== Final fit on full data: log_reg ====
Best params: {'clf__C': 0.1}
Mean inner ROC AUC: 0.960

==== Final fit on full data: svm_rbf ====
Best params: {'clf__C': 1.0, 'clf__gamma': 0.01}
Mean inner ROC AUC: 0.946

==== Final fit on full data: gp_classifier ====
Best params: {'clf__kernel': RBF(length_scale=10)}
Mean inner ROC AUC: 0.949
```



## 6.2 Decision regions in a shared 2D feature plane

From the permutation importance plots, all three models consistently rank


- `R_hf_ohm` (high-frequency / ohmic resistance), and
- `Zmag_min` (minimum impedance magnitude across the spectrum)

as the two most informative features.

To make the decision regions directly comparable, I now:

- fix a **common 2D feature plane**  
(`R_hf_ohm` vs `Zmag_min`),

- train a 2D version of each final model on these two features only, and
- plot the resulting **probability maps** and data points in that plane.



In this shared ( $R_{hf\_ohm}$ ,  $Zmag\_min$ ) plane, all three models agree that cells with **higher ohmic resistance and higher minimum impedance** sit in a region of high probability of being *aged* (yellow, top-right), while low values on both axes are confidently *healthy* (purple, bottom-left). The main difference is how sharply each model transitions through the **overlap band** where healthy and aged points mix: logistic regression uses an almost straight boundary, whereas the SVM and GP give slightly smoother transitions, but they all recover essentially the same separation structure.